

SUMMARY

Data Engineer with 5+ years of experience designing scalable data pipelines and analytics solutions in insurance, healthcare, sports analytics and AI/ML domains. Skilled in Python, PySpark, Spark, Kafka, Airflow, AWS, and Azure, with expertise in real-time and batch data processing at scale. Proven track record of reducing pipeline runtimes by up to 86%, optimizing data architectures, and deploying production ML systems that drive measurable business impact..

CERTIFICATIONS & AWARDS

AWS Certified Data Engineer – Associate (DEA-C01) | Microsoft Certified: Fabric Analytics Engineer Associate (DP-700) | DEVHACKS 2026 – 1st Place (100+ teams) | SPOT Recognition Award – Super Six Sports Gaming

SKILLS

- Programming & Scripting: Python, SQL, PySpark
- Big Data Technologies: Spark Structured Streaming, Apache Kafka, Apache Hive, Hadoop, MapReduce, Apache Airflow (MWAA)
- Cloud Platforms & Services: AWS (Glue, EMR, Redshift, S3, Athena, Kinesis, Lambda, MWAA, CloudWatch), Azure (Databricks, ADF), Snowflake, Microsoft Fabric
- Databases & Storage: Snowflake, PostgreSQL, MySQL, MS SQL Server, MongoDB
- AI/ML & LLM Systems: LangChain, LangGraph, Agentic RAG, Vector Databases, Claude/GPT-4 APIs, Scikit-Learn, PyTorch
- DevOps & Infrastructure: Docker, Terraform, Jenkins, Git, CI/CD, Data Security (IAM, RBAC)

EXPERIENCE

AI Engineer, MyEdMaster – Industry Capstone Project (ASU SER517)

Jan 2026 – Apr 2026

- Developed multi-agent Stateful RAG system for personalized legal guidance using LangGraph with Qdrant vector DB, achieving sub-2s query latency across 10K+ documents and 95% relevance accuracy.
- Built FastAPI + Node.js backend with Docker supporting 100+ concurrent sessions with sub-200ms response time through optimized vector retrieval.

Data Engineer II, EXL Services

Jul 2023 – Mar 2024

- Drove 40% BI query performance improvement by architecting S3 data lake and Snowflake warehouse across 8 datasets supporting 500K+ daily users processing high-frequency insurance claim events.
- Eliminated 90-minute production bottleneck by optimizing PySpark ETL pipelines on AWS Glue/MWAA processing 100M records/batch, achieving 66% runtime reduction via partition pruning and Spark query optimization.
- Built production anomaly detection layer using Great Expectations across 15+ data sources with CloudWatch monitoring and MWAA retries, maintaining 100% SLAs.

Data Engineer, Super Six Sports Gaming

Aug 2022 – Jul 2023

- Reduced user churn 20% by building production ML retention pipeline with 78% accuracy using Scikit-Learn and statistical modeling that identified behavioral anomalies in time-series user activity data, deployed automated Python/SQL re-engagement campaigns.
- Built and owned core data ingestion layer — multi-source batch ETL pipelines ingesting from 10+ REST APIs and S3 using cloud data solutions, loading 500K+ daily records of football, soccer, and basketball data into MongoDB with sub-hourly refresh cadence, enabling near-real-time analytics for product and marketing teams with 99.9% data accuracy via cross-source validation and data profiling
- Engineered ML feature pipelines with SCD Type 1/2 dimensional modeling and automated feature engineering, reducing ML experiment cycle time 35% through optimized time-series data storage patterns.

Associate Data Engineer, Futureense Technologies

Oct 2021 – Jul 2022

- Spearheaded migration of 1 Billion+ record oncology pipelines from legacy SAS to Apache Spark on Azure Databricks, achieving 86% reduction in batch processing time from 6+ hours to 50 minutes through broadcast joins and strategic partitioning.
- Automated bi-weekly HCP targeting reports by building Python ETL pipeline querying AWS Athena, reducing manual effort from 4–5 hours to 15 minutes, delivering 95% time savings.
- Delivered 99.5% data accuracy post-migration via comprehensive PySpark and SQL validation frameworks with statistical reconciliation.

Data Analyst, Koron Projects Limited

Oct 2018 – Jul 2021

- Built and maintained 15 Power BI dashboards with advanced analytics for executive leadership tracking project costs, timelines, and profitability across \$50M+ in annual construction projects.
- Consolidated cost data from multiple enterprise sources including SQL Server, Oracle, and MySQL via SQL queries and stored procedures, automating monthly executive reporting and reducing manual effort by 20 hours/month.

EDUCATION

Master in Software Engineering (AI Specialization), Arizona State University

Aug 2024 – May 2026

Bachelor in Computer Science, IP University — New Delhi, India

Aug 2014 – Jun 2018