

HIMANSHU PANT

Tempe, AZ (Open to Relocation & Remote) | 623-999-4043 | hpant.data@gmail.com | [LinkedIn](#) | [Github](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Data Engineer with 3+ years of experience building scalable ETL/ELT pipelines, real-time data infrastructure, and cloud-native data platforms across insurance, sports analytics, and healthcare. Proven track record of processing 100M+ records per batch and migrating datasets exceeding 1B records to distributed cloud architectures, reducing pipeline runtimes by up to 86%. Expert in data warehousing solutions, data lake management, and big data technologies including PySpark, Airflow (MWAA), DBT, Kafka, AWS, Databricks, Snowflake, and Microsoft Fabric.

TECHNICAL SKILLS

Languages	Python, SQL, PySpark, Scala
Data Engineering	Apache Spark, Spark Structured Streaming, Apache Kafka, Apache Airflow (MWAA), DBT, Apache Hive, Hadoop, ETL/ELT, Data Warehousing, Dimensional Modeling (SCD Type 1/2), Data Profiling, Data Quality Tools (Great Expectations), Data Lake Management, Big Data Technologies
Cloud & Infra	AWS (S3, Glue, EMR, Redshift, Kinesis, Firehose, Lambda, Athena, IAM, CloudWatch, MWAA), Azure (Databricks, ADF), Microsoft Fabric, Cloud Data Solutions, Data Storage Solutions, Terraform (IaC), Docker, Jenkins, CI/CD
Databases	Snowflake, AWS Redshift, PostgreSQL, MSSQL, MongoDB, SQLite, Data Warehousing Solutions
AI / ML	LangChain, LangGraph, Agentic RAG, ChromaDB, Vector Databases, Claude/GPT-4 API, SapBERT
Analytics & BI	Power BI, Pandas, NumPy, Matplotlib, Data Analysis
Security & Governance	Data Security (IAM, RBAC), Data Governance, Role-Based Access Control

PROFESSIONAL EXPERIENCE

AI Engineer <i>MyEdMaster - Capstone Project</i>	Jan 2026 – Apr 2026 <i>Remote</i>
- Built multi-agent RAG system for personalized legal guidance — created a LangGraph orchestration pipeline with Qdrant vector DB, achieving sub-2s query latency across 10K+ legal documents using advanced data profiling and quality validation techniques - Designed data ingestion and transformation pipelines converting unstructured legal corpora into structured, queryable vector embeddings with real-time personalization signals via LLM APIs, achieving 95% relevance accuracy for personalized legal recommendations - Built FastAPI + Node.js backend integrating vector DB lookups, session state management, and multi-turn conversation history, supporting 100+ concurrent sessions with < 200ms response time for verified legal benchmarks	
Data Engineer II <i>EXL Services</i>	Jul 2023 – Mar 2024 <i>Gurgaon, India</i>
- Eliminated 90-minute production bottleneck in large-scale insurance claims processing by redesigning PySpark ETL pipelines on AWS Glue and MWAA — processing ~100M records per batch (~50–150 GB) with 66% runtime reduction via partition pruning, predicate pushdown, and query plan optimization leveraging big data technologies - Architected S3-based data lake with Snowflake data warehousing solution across 8 datasets, implementing dual-sink architecture (Snowflake for BI consumption, S3 for downstream ML feature pipelines) with comprehensive data lake management practices — reducing BI query time by 40% and enabling parallel ML feature engineering without impacting production analytics - Built production observability layer with MWAA retries, email alerting, and CloudWatch dashboards; enforced data quality across 15+ upstream insurance feeds using Great Expectations and data quality tools, ensuring 100% validity SLAs pre-ingestion with automated data profiling and anomaly detection - Engineered Airflow DAGs with parallel and dependent task chains; accelerated deployment velocity by 40% via Jenkins CI/CD pipelines automating Spark job delivery to AWS Glue and EMR across a 3-person engineering pod - Standardized cloud data solutions and data storage infrastructure by authoring Terraform IaC modules enforcing data security via role-based IAM access controls, eliminating environment drift across dev/staging/prod	
Data Engineer <i>Super Six Sports Gaming</i>	Aug 2022 – Jul 2023 <i>Gurgaon, India</i>

- **Built and owned** core data ingestion layer — multi-source batch ETL pipelines ingesting from 10+ REST APIs and S3 using cloud data solutions, loading 500K+ daily records of football, soccer, and basketball data into MongoDB with sub-hourly refresh cadence, enabling near-real-time analytics for product and marketing teams with **99.9% data accuracy** via cross-source validation and data profiling
- Engineered high-performance ML feature pipelines with SCD Type 1/2 dimensional modeling and data warehousing best practices — automating feature engineering and reducing ML experiment cycle time by **35%** through optimized data storage solutions
- Drove **20%** reduction in user churn by building production ML retention pipeline (**78% accuracy**) — translating raw behavioral data into automated Python/SQL-based re-engagement campaigns targeting at-risk segments using data analysis and profiling techniques
- Received **SPOT Recognition Award** for Ownership and outstanding delivery of various ETL frameworks under tight deadlines demonstrating cross-functional collaboration on real-time data infrastructure build-out

Associate Data Engineer
Futureense Technologies

Oct 2021 – Jul 2022
Bengaluru, India

- Led migration of **1B+** record oncology patient pipelines from legacy SAS batch jobs to Apache Spark on Azure Databricks — ingesting from S3 and writing curated outputs leveraging big data technologies and cloud data solutions, converting manual workflows into fully automated, distributed cloud-native processes that enabled 24/7 processing without manual intervention
- Achieved **86%** reduction in critical batch processing time (6+ hours → under 50 minutes) by implementing broadcast joins, strategic partitioning, and data skew resolution in PySpark with advanced data profiling to identify bottlenecks
- Automated bi-weekly HCP targeting report generation by building Python ETL pipeline (Boto3 + Pandas) to query AWS Athena and generate formatted Excel reports using data analysis techniques, reducing manual effort from 4-5 hours to 15 minutes per cycle (**95% time savings**) and eliminating data extraction errors
- Delivered 99.5% data accuracy post-migration by developing comprehensive PySpark and SQL-based validation frameworks with data quality tools, maintaining production data quality standards across all migrated datasets

PROJECTS

MeetFlow | github.com/hpant5/MeetFlow

Mar 2026

- Designed two-agent LLM architecture (transcript analyzer + capacity-aware smart assigner) that parses meeting transcripts, extracts actionable tasks via GPT-4o, and auto-assigns them to Jira/Taiga based on real-time team workload — won Track 1 at DEVHACKS 2026 competing against 100+ teams

MakeLifeEasy (Daybreak) | LA Hacks 2026

Apr 2026

- Built multi-agent AI orchestration system on Fetch.ai's AgentVerse platform coordinating task execution across Gmail, Google Calendar, Notion, Jira, and GitHub via natural language — architected 4 specialized uAgents leveraging LangGraph workflow orchestration and Claude API for intent classification, eliminating context-switching and enabling single-prompt multi-platform automation

FairCharge | Claude AI Builder Club Hackathon

Mar 2026

- Built end-to-end medical bill audit pipeline: Claude Vision parses bills via OCR, hybrid matcher (SQLite exact match + SapBERT semantic + ChromaDB fuzzy) benchmarks charges against 19,226 CMS Medicare CPT codes, rule engine auto-generates dispute letters — handling medical abbreviations and OCR errors with 92% accuracy

EDUCATION

Arizona State University
M.S., Software Engineering, Data Science Minor

Aug 2024 – May 2026
Tempe, AZ

IP University
B.Tech., Computer Science

Aug 2014 – Jun 2018
New Delhi

CERTIFICATIONS & AWARDS

Microsoft Certified: Fabric Analytics Engineer Associate (DP-700): May 2026

HackerRank: SQL Certified

DEVHACKS 2026 - 1st Place, Track 1: MeetFlow - AI Meeting Orchestration (100+ Teams)

SPOT Recognition Award: Super Six Sports Gaming - data infrastructure excellence