

Himanshu Pant

Data Engineer

Tempe, AZ | 623-999-4043 | hpant.data@gmail.com

linkedin.com/in/himanshupant-de | github.com/hpant5 | himanshupant.dev

SUMMARY

Data engineer building production-grade distributed pipelines and agentic AI systems. Specialized in **Apache Spark optimization**, **AWS-native data infrastructure**, and **LangGraph-based LLM applications**, with a track record of cutting batch runtimes by **86% on 1B+ record pipelines** and shipping data and ML systems serving **500K+ daily users** across insurance, healthcare, and sports analytics. **AWS and Microsoft certified**; comfortable owning systems end-to-end from ingestion through analytics and ML deployment.

EDUCATION

Arizona State University

Master of Science in Software Engineering (AI Specialization)

Tempe, AZ

Aug 2024 – May 2026

IP University

Bachelor of Technology in Computer Science

New Delhi, India

Aug 2014 – Jun 2018

EXPERIENCE

Data & AI Engineer (Industry Capstone)

Jan 2026 – Apr 2026

MyEdMaster

Tempe, AZ

- Developed multi-agent stateful RAG system for personalized legal guidance using **LangGraph** and **Qdrant vector DB**, achieving **sub-2s query latency** across 10K+ documents and **95% relevance accuracy**.
- Built **FastAPI + Node.js** backend with **Docker** supporting 100+ concurrent sessions at **sub-200ms response time** via optimized retrieval and a 5-node pipeline.
- Authored evaluation framework for the QnA service and resolved critical production **NameError** in the agentic graph execution layer; shipped **4-layer technical documentation** adopted by partner team.

Consultant II, Data Analytics & Engineering

Jul 2023 – Mar 2024

EXL Services

Gurugram, India

- Eliminated a **90-minute production bottleneck** by optimizing **PySpark ETL on AWS Glue / MWAA** processing 100M+ records per batch, cutting runtimes **66% via partition pruning**.
- Architected **S3 data lake and Snowflake warehouse** across 8 datasets for **500K+ daily users** on high-frequency insurance claim events, driving **40% BI query performance improvement**.
- Built production data quality framework using **Great Expectations** across 15+ data sources with **CloudWatch monitoring**, maintaining **100% SLA compliance**.
- Automated infrastructure provisioning by deploying Airflow DAGs via **Terraform IaC** and integrating **Jenkins CI/CD** with unit and integration testing.

Data Engineer

Aug 2022 – Jul 2023

Super Six Sports Gaming

Gurugram, India

- Built and owned the core data ingestion layer: multi-source **batch ETL from 10+ REST APIs and S3**, loading **500K+ daily records** of sports event data into **MongoDB** with **99.9% accuracy**.
- Reduced user churn 20%** by building a production **ML retention pipeline (78% accuracy, Scikit-Learn)** detecting behavioral anomalies in time-series user activity and triggering automated Python/SQL workflows.
- Engineered ML feature pipelines with **SCD Type 1/2 dimensional modeling**, accelerating ML experiment cycle time by **35%**.

Associate Data Engineer

Oct 2021 – Jul 2022

Futureense Technologies

Bangalore, India

- Spearheaded migration of **1B+ record oncology pipelines** from legacy **SAS** to **Apache Spark on Azure Databricks**, cutting batch time from 6+ hours to 50 minutes (**86% reduction**) via broadcast joins.
- Automated bi-weekly HCP targeting reports via **Python ETL on AWS Athena**, reducing manual effort from 4-5 hours to 15 minutes (**95% savings**).
- Delivered **99.5% data accuracy** post-migration via comprehensive **PySpark + SQL validation frameworks** with statistical reconciliation.

Data Analyst

Oct 2018 – Jul 2021

Koron Projects Limited

New Delhi, India

- Built 15 **Power BI dashboards** with advanced analytics for executive leadership tracking project costs, timelines, and profitability across **\$50M+ in annual construction projects**.
- Consolidated cost data from SQL Server, Oracle, and MySQL via **SQL and stored procedures**, removing **20 hours/month of manual effort**.

PROJECTS

- MeetFlow** | *Python, LangGraph, Claude API, Jira/Taiga, Slack* | **DEVHACKS 2026 Winner** Mar 2026
- Built a **multi-agent meeting intelligence system** that ingests transcripts, queries Jira / Taiga for current engineer workload, and generates LLM-driven task assignment recommendations balanced against capacity.
 - Implemented **human-in-the-loop approval via Slack**: admins receive recommendation plus rationale and approve before tasks are written back to tracker systems.
- Daybreak (MakeLifeEasy)** | *Fetch.ai Agentverse, LangGraph, REST integrations* | **LA Hacks 2026** Apr 2026
- Built **multi-agent day planner** on **Fetch.ai Agentverse** using **LangGraph** and direct REST integrations across Gmail, Google Calendar, Notion, Jira, and GitHub.
 - Specialized uAgents handle intent parsing, calendar reasoning, and cross-system task coordination, returning a synthesized daily plan from natural language input.
- JobGenie** | *FastAPI, Playwright, Supabase, Netlify, Render* | jobgenie.dev 2026
- Production **job-application pipeline** ingesting from Greenhouse, Lever, and Adzuna via **FastAPI + Playwright scrapers**, persisted to **Supabase** with deduplication.
 - Multi-resume upload and per-application tracking; deployed live on **Netlify + Render** with active usage.

TECHNICAL SKILLS

Languages: Python, SQL, PySpark, Scala

Distributed & Big Data: Apache Spark (Structured Streaming), AWS Glue, Azure Databricks, Apache Kafka, Hive, Hadoop, dbt

Cloud & Data Platforms: AWS (Glue, EMR, Redshift, S3, Athena, Kinesis, Lambda, MWAA, CloudWatch), Azure (Databricks, ADF), Snowflake, Microsoft Fabric

Storage & Databases: Snowflake, Amazon Redshift, PostgreSQL, MySQL, MS SQL Server, MongoDB, Qdrant, pgvector

Orchestration & DevOps: Apache Airflow (MWAA), Docker, Terraform, Jenkins, Git, CI/CD, Great Expectations

AI / ML & LLM Systems: LangChain, LangGraph, Agentic RAG, Vector Databases, Claude/GPT-4 APIs, PyTorch, Scikit-Learn, FastAPI

CERTIFICATIONS & AWARDS

- **AWS Certified Data Engineer — Associate (DEA-C01)**
- **Microsoft Certified: Fabric Analytics Engineer Associate (DP-700)**
- **DEVHACKS 2026 — 1st Place (100+ teams)** for MeetFlow
- **SPOT Recognition Award** — Super Six Sports Gaming

Work Authorization: Eligible to work in the US through ~2029 (F-1 OPT + STEM Extension; no immediate sponsorship required).