

# HIMANSHU PANT

Tempe, AZ (Open to Relocation & Remote) [hpant.data@gmail.com](mailto:hpant.data@gmail.com) | [LinkedIn](#) | [Github](#) | [Portfolio](#)

## PROFESSIONAL SUMMARY

Data Engineer with **5+ years** building scalable ETL/ELT pipelines, real-time data infrastructure, and AI-enabled analytics platforms. Delivered measurable product impact: reduced user churn **20%**, accelerated BI query performance **40%**, and processed **100M+ records/batch** with **66%** runtime improvements. Expert in translating complex data into actionable insights for product, engineering, and business teams using PySpark, Airflow (MWAA), Kafka, Snowflake, and AWS. Recently expanded into Agentic RAG and LLM orchestration (M.S. Software Engineering, ASU, May 2026).

## TECHNICAL SKILLS

|                        |                                                                                                                                                                                                                                                                         |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Languages              | Python, SQL, PySpark, Scala                                                                                                                                                                                                                                             |
| Cloud & Infrastructure | AWS (Glue, EMR, Redshift, S3, Athena, Kinesis, Lambda, MWAA, CloudWatch), Azure (Databricks, ADF), Snowflake, Microsoft Fabric, Docker, Terraform, CI/CD (Jenkins)                                                                                                      |
| Data Engineering       | Apache Spark, Spark Structured Streaming, Apache Kafka, Apache Airflow (MWAA), DBT, Apache Hive, Hadoop, ETL/ELT, Data Warehousing, Dimensional Modeling (SCD Type 1/2), Data Profiling, Data Quality (Great Expectations), Data Lake Management, Big Data Technologies |
| Databases              | Snowflake, AWS Redshift, PostgreSQL, MSSQL, MongoDB, SQLite                                                                                                                                                                                                             |
| AI / ML                | LangChain, LangGraph, Agentic RAG, ChromaDB, Vector Databases, Claude/GPT-4 API, SapBERT                                                                                                                                                                                |
| Analytics & BI         | Power BI, Pandas, NumPy, Matplotlib, Data Analysis                                                                                                                                                                                                                      |

## PROFESSIONAL EXPERIENCE

|                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |                                              |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------|
| <b>AI Engineer</b><br><i>MyEdMaster - Capstone Project</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | Jan 2026 – Apr 2026<br><i>Remote</i>         |
| <ul style="list-style-type: none"><li>Developed multi-agent RAG system for legal guidance using LangGraph orchestration with Qdrant vector DB, achieving <b>sub-2s</b> query latency across <b>10K+</b> documents via advanced data profiling and validation</li><li>Designed data pipelines converting unstructured legal corpora into queryable vector embeddings with real-time personalization via LLM APIs, achieving <b>95%</b> relevance accuracy</li><li>Built FastAPI + Node.js backend with vector DB integration, session management, and conversation history supporting <b>100+</b> concurrent sessions with <b>&lt; 200 ms</b> response time</li></ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                                              |
| <b>Data Engineer II</b><br><i>EXL Services</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | Jul 2023 – Mar 2024<br><i>Gurgaon, India</i> |
| <ul style="list-style-type: none"><li>Drove <b>40%</b> BI query performance improvement for product and analytics teams by architecting S3 data lake + Snowflake warehouse across <b>8 datasets</b> with dual-sink pattern (Snowflake for BI consumption, S3 for ML features) — enabled parallel ML experimentation without production impact while supporting <b>500K+</b> daily active insurance users</li><li>Eliminated <b>90-minute</b> production bottleneck in claims processing by redesigning PySpark ETL pipelines on AWS Glue/MWAA processing <b>~100M records/batch (50–150 GB)</b>, achieving <b>66%</b> runtime reduction via partition pruning, predicate pushdown, and query optimization</li><li>Built production observability layer with CloudWatch dashboards tracking pipeline health across <b>15+ data sources</b>, MWAA retries, and email alerting; enforced data quality using Great Expectations ensuring <b>100%</b> validity SLAs with automated anomaly detection</li><li>Accelerated deployment velocity <b>40%</b> by engineering Airflow DAGs with parallel task chains and building Jenkins CI/CD pipelines automating Spark job delivery to AWS Glue/EMR across <b>3-person</b> engineering team</li><li>Standardized infrastructure using Terraform IaC modules enforcing role-based IAM access controls, eliminating environment drift across dev/staging/prod</li></ul> |                                              |
| <b>Data Engineer</b><br><i>Super Six Sports Gaming</i>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | Aug 2022 – Jul 2023<br><i>Gurgaon, India</i> |
| <ul style="list-style-type: none"><li>Reduced user churn <b>20%</b> by partnering with product team to build production ML retention pipeline (<b>78%</b> accuracy) that identified at-risk user segments and translated raw behavioral data into automated Python/SQL re-engagement campaigns — directly influenced product roadmap for retention features</li><li>Built and owned core data ingestion layer — multi-source ETL pipelines from <b>10+ REST APIs</b> and S3, loading <b>500K+</b> daily records into MongoDB with sub-hourly refresh cadence, enabling near-real-time analytics for product and marketing teams with <b>99.9%</b> accuracy</li></ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                                              |

- Engineered ML feature pipelines with SCD Type 1/2 dimensional modeling, automating feature engineering and reducing ML experiment cycle time **35%** through optimized data storage patterns
- Received SPOT Recognition Award for ownership and outstanding delivery of data infrastructure under tight deadlines

### Associate Data Engineer

*Futureense Technologies*

Oct 2021 – Jul 2022

*Bengaluru, India*

- Spearheaded migration of **1B+ record** oncology pipelines from legacy SAS to Apache Spark on Azure Databricks, converting manual workflows into fully automated cloud-native processes enabling 24/7 processing
- Achieved **86%** reduction in batch processing time (**6+ hours** → **50 minutes**) by implementing broadcast joins, strategic partitioning, and data skew resolution with advanced profiling
- Automated bi-weekly HCP targeting reports by building Python ETL pipeline (Boto3 + Pandas) querying AWS Athena, reducing manual effort from **4-5 hours** to **15 minutes** (**95%** time savings) and eliminating extraction errors
- Delivered **99.5%** data accuracy post-migration via comprehensive PySpark and SQL validation frameworks, maintaining production quality standards across all datasets

### Data Analyst

*Koron Projects Limited*

Oct 2018 – Jul 2021

*Gurgaon, India*

- Built and maintained **15 Power BI dashboards** for executive leadership tracking project costs, timelines, and profitability across **\$50M+** in annual construction projects — enabled data-driven decision-making for **50+** active projects
- Consolidated cost data from multiple enterprise sources (SQL Server, Oracle) via SQL queries and stored procedures, automating monthly executive reporting and reducing manual effort by **20 hours/month** (replacing Excel-based processes)
- Collaborated with cross-functional project teams to translate business requirements into KPI dashboards, identifying budget variances and informing strategic decisions across finance and operations stakeholders

## PROJECTS

**MeetFlow** | [github.com/hpant5/MeetFlow](https://github.com/hpant5/MeetFlow)

Mar 2026

- Designed two-agent LLM architecture (transcript analyzer + capacity-aware assigner) parsing meeting transcripts, extracting tasks via GPT-4o, and auto-assigning to Jira/Taiga based on team workload — won **1st Place** at DEVHACKS 2026 (**100+** teams)

**MakeLifeEasy (Daybreak)** | LA Hacks 2026

Apr 2026

- Built multi-agent orchestration system on Fetch.ai's AgentVerse coordinating tasks across Gmail, Google Calendar, Notion, Jira, and GitHub via natural language — architected **4 specialized** uAgents using LangGraph workflow orchestration and Claude API for intent classification

**FairCharge** | Claude AI Builder Club Hackathon

Mar 2026

- Built medical bill audit pipeline: Claude Vision OCR + hybrid matcher (SQLite exact + SapBERT semantic + ChromaDB fuzzy) benchmarks charges against **19,226** CMS Medicare CPT codes, auto-generating dispute letters with **92%** accuracy

## EDUCATION

**Arizona State University**

*M.S., Software Engineering (AI Specialization), Data Science Minor*

Aug 2024 – May 2026

*Tempe, AZ*

**IP University**

*B.Tech., Computer Science*

Aug 2014 – Jun 2018

*New Delhi*

## CERTIFICATIONS & RECENT AWARDS

**Microsoft Certified:** Fabric Analytics Engineer Associate (DP-700) – May 2026

**DEVHACKS 2026** – 1st Place, Track 1 (100+ teams)

**SPOT Recognition Award** – Super Six Sports Gaming

**HackerRank** – SQL Certified