

Himanshu Pant

Data Engineer

Tempe, AZ | 623-999-4043 | hpant.data@gmail.com | linkedin.com/in/himanshupant-de | himanshupant.dev

PROFESSIONAL SUMMARY

AWS- and Microsoft-certified **Data Engineer** with 5+ years designing and operating **large-scale batch and streaming pipelines** across insurance, sports analytics, and healthcare. Hands-on with **Python, SQL, PySpark, dbt, Airflow**, and the **AWS and Azure Databricks** ecosystems, with growing depth in **AI/LLM systems** (RAG, LangGraph). Track record of cutting pipeline runtimes 66%, migrating 1B+ record workloads to the cloud, and mentoring junior engineers. M.S. in Software Engineering from Arizona State University.

EDUCATION

Master of Science in Software Engineering (Data Science Minor), Arizona State University

May 2026

Bachelor of Technology in Computer Science, IP University — New Delhi, India

August 2014 – June 2018

PROFESSIONAL EXPERIENCE

AI / Data Engineer (Industry Capstone) – *MyEdMaster*

Tempe, AZ | January 2026 – April 2026

- Built a stateful, multi-agent **RAG** system using **LangGraph** and a **Qdrant** vector DB, achieving sub-2s query latency across 10K+ documents.
- Designed ingestion and transformation pipelines converting unstructured corpora into structured, queryable vector embeddings with real-time personalization signals via **LLM APIs**.
- Engineered a **FastAPI** + Node.js backend (Dockerized) with an optimized 5-node retrieval pipeline; authored the evaluation framework and layered technical documentation adopted by the partner team.

Consultant II, Data Engineering – *EXL Services*

Gurugram, India | July 2023 – March 2024

- Delivered across two client engagements (a major US auto insurer and a global toy manufacturer); eliminated a 90-minute production bottleneck by optimizing **PySpark ETL** on **AWS Glue** + MWAA processing 100M+ records/batch, cutting runtimes **66%** via partition pruning and predicate pushdown.
- Architected an **S3** data lake and **Snowflake** warehouse with a **Great Expectations** data-quality framework across 15+ sources at 100% SLA compliance; built an Apache Iceberg feature store for fraud-detection ML models (Glue, S3, SageMaker).
- Architected end-to-end ETL integrating ADLS, SharePoint, APIs, and flat files into structured mart layers using advanced **dbt** (Jinja macros, snapshots for SCD, incremental models); shipped CI/CD via **Jenkins** and dbt docs with column-level lineage. Mentored 2 junior engineers.

Data Engineer – *Super Six Sports Gaming*

Gurugram, India | August 2022 – July 2023

- Owned core data ingestion layers, building multi-source batch **ETL** pipelines from 10+ REST APIs into MongoDB, with automated validation and redundancy checks that eliminated data-integrity errors in routine spot checks.
- Built a production ML user-retention pipeline (**78%** accuracy, validated against a control group) detecting behavioral anomalies in user-activity data, driving a **20%** reduction in churn.
- Engineered ML feature pipelines with SCD Type 1/2 dimensional modeling; promoted within 6 months and awarded SPOT Recognition for ownership and delivery.

Associate Data Engineer – *Futureense Technologies*

Bangalore, India | October 2021 – July 2022

- Co-led the migration of 1B+ oncology records from legacy SAS to **Spark** on **Azure Databricks**, using partitioning and broadcast joins to cut batch execution from 6 hours to 50 minutes.
- Automated a previously manual recurring healthcare reporting pipeline with a parameterized Python **ETL** layer on **AWS Athena**.
- Validated the migration with comprehensive **PySpark** and **SQL** reconciliation frameworks; zero data-integrity issues reported post-migration, with supporting dataset documentation and metadata.

Executive Analyst, Data & Reporting – *Koron Projects Limited*

New Delhi, India | October 2018 – July 2021

- Built and maintained **ETL** pipelines extracting and cleaning enterprise cost data from **Oracle**, MS SQL Server, and MySQL via **SQL** stored procedures, feeding downstream BI reporting across a \$50M project portfolio.
- Designed **dimensional reporting models** and built 15+ **Power BI** dashboards giving leadership real-time visibility into project costs, timelines, and profitability across 50+ sites; automated recurring reporting, saving ~20 hours/month.

SKILLS

Languages & Core: Python, SQL, PySpark, ETL/ELT, Dimensional Data Modeling, OLAP, Spark Structured Streaming, dbt

Databases & Warehousing: Snowflake, Redshift, Oracle, MS SQL Server, MySQL, MongoDB, Apache Iceberg, Kafka, Hive, Hadoop

AI & Cloud: LangGraph, RAG, Qdrant, FastAPI, LLM APIs, AWS (Glue, Lambda, S3, Redshift, Athena, EMR, Kinesis, IAM), Azure Databricks

Orchestration, BI & Tools: Airflow (MWAA), Great Expectations, Terraform, Power BI, Docker, Jenkins CI/CD, Git, Agile/Scrum

CERTIFICATIONS

AWS Certified Data Engineer – Associate | Microsoft Certified: Fabric Data Engineer Associate