

Himanshu Pant — Data Engineer

623-999-4043 | hpant.data@gmail.com | linkedin.com/in/himanshupant-de | himanshupant.dev

SUMMARY

Data Engineer with 5+ years of expertise designing scalable data pipelines and analytics solutions across insurance, sports analytics, healthcare, and AI domains. Skilled in Python, PySpark, Spark, Kafka, Airflow, and cloud platforms (AWS, Azure) to enable real-time and batch processing at scale. Proven ability to optimize data architecture, reduce pipeline runtimes by up to 86%, and deploy production ML systems that deliver measurable business impact through robust engineering practices and cross-functional collaboration.

CERTIFICATIONS & AWARDS

AWS Certified Data Engineer – Associate (DEA-C01) | Microsoft Certified: Fabric Analytics Engineer Associate (DP-700) | DEVHACKS 2026 – 1st Place (100+ teams) | SPOT Recognition Award – Super Six Sports Gaming

SKILLS

Programming & Scripting: Python, SQL, PySpark

Big Data Technologies: Spark Structured Streaming, Apache Kafka, Apache Hive, Hadoop, MapReduce, Apache Airflow (MWAA)

Cloud Platforms & Services: AWS (Glue, EMR, Redshift, S3, Athena, Kinesis, Lambda, MWAA, CloudWatch), Azure (Databricks, ADF), Snowflake, Microsoft Fabric

Databases & Storage: Snowflake, PostgreSQL, MySQL, MS SQL Server, MongoDB

AI/ML & LLM Systems: LangChain, LangGraph, Agentic RAG, Vector Databases (Qdrant, ChromaDB), Claude/GPT-4 APIs, Scikit-Learn, PyTorch

DevOps & Infrastructure: Docker, Terraform, Jenkins, Git, CI/CD, Data Security (IAM, RBAC)

EXPERIENCE

AI Engineer, MyEdMaster – Industry Capstone Project (ASU SER517)

Jan 2026 – Apr 2026

- Developed multi-agent RAG system for personalized legal guidance using LangGraph with Qdrant vector DB, achieving sub-2s query latency across 10K+ documents and 95% relevance accuracy.
- Built FastAPI + Node.js backend with Docker supporting 100+ concurrent sessions with sub-200ms response time through optimized vector retrieval.

Data Engineer II, EXL Services

Jul 2023 – Mar 2024

- Drove 40% BI query performance improvement by architecting S3 data lake and Snowflake warehouse across 8 datasets supporting 500K+ daily users processing high-frequency insurance claim events.
- Eliminated 90-minute production bottleneck by redesigning PySpark ETL pipelines on AWS Glue/MWAA processing 100M records/batch, achieving 66% runtime reduction via partition pruning and Spark query optimization.
- Built production anomaly detection layer using Great Expectations across 15+ data sources with CloudWatch monitoring and MWAA retries, maintaining 100% SLAs.

Data Engineer, Super Six Sports Gaming

Aug 2022 – Jul 2023

- Reduced user churn 20% by building production ML retention pipeline with 78% accuracy using Scikit-Learn and statistical modeling that identified behavioral anomalies in time-series user activity data, deployed automated Python/SQL re-engagement campaigns.
- Built and owned real-time data ingestion layer using multi-source ETL pipelines from 10+ REST APIs and S3 with Kafka-based streaming, loading 500K+ daily records into MongoDB with sub-hourly refresh cadence, enabling near-real-time analytics with 99.9% accuracy.
- Engineered ML feature pipelines with SCD Type 1/2 dimensional modeling and automated feature engineering, reducing ML experiment cycle time 35% through optimized time-series data storage patterns.

Associate Data Engineer, Futureense Technologies

Oct 2021 – Jul 2022

- Spearheaded migration of 1B+ record oncology pipelines from legacy SAS to Apache Spark on Azure Databricks, achieving 86% reduction in batch processing time from 6+ hours to 50 minutes through broadcast joins and strategic partitioning.
- Automated bi-weekly HCP targeting reports by building Python ETL pipeline querying AWS Athena, reducing manual effort from 4–5 hours to 15 minutes, delivering 95% time savings.
- Delivered 99.5% data accuracy post-migration via comprehensive PySpark and SQL validation frameworks with statistical reconciliation.

Data Analyst, Koron Projects Limited

Oct 2018 – Jul 2021

- Built and maintained 15 Power BI dashboards with advanced analytics for executive leadership tracking project costs, timelines, and profitability across \$50M+ in annual construction projects.
- Consolidated cost data from multiple enterprise sources including SQL Server, Oracle, and MySQL via SQL queries and stored procedures, automating monthly executive reporting and reducing manual effort by 20 hours/month.

EDUCATION

Master in Software Engineering (AI Specialization), Arizona State University — Tempe, AZ

Aug 2024 – May 2026

Bachelor in Computer Science, IP University — New Delhi, India

Aug 2014 – Jun 2018